

Admin User Guide

How to run, read, and get the most out of your web-push notification system.

This guide is written for editors and administrators — no technical background needed. It explains, in plain language, what every screen does, what the numbers mean, what to expect when you see certain values, and exactly how to send a notification yourself.

The 30-second version

The system reads Taxscan's articles automatically and pushes only the important ones to your readers' browsers, paced sensibly so nobody is annoyed. Your job is light: **glance at the Dashboard** to check things are healthy, and **clear the Review queue** (decide on the few articles the system wasn't sure about). You can also **send anything yourself, instantly**, whenever you want.

What's inside

- 1 What this tool is (in one minute)
- 2 Logging in & your account
- 3 The big picture: how notifications go out
- 4 Your daily routine — the 5-minute check
- 5 Reading the Dashboard (what every number means)
- 6 The Review queue — your main task
- 7 Sending a notification yourself (Compose)
- 8 The other screens — Campaigns, Reports, Reads & Surfaces
- 9 Quick reference card
- 10 Frequently asked questions

1 What this tool is (in one minute)

Taxscan Push sends short notifications straight to your readers' web browsers — even when they aren't on the Taxscan website — to tell them about important new updates. A reader "subscribes" once by clicking *Allow* on a prompt; after that, the notifications appear on their device.

The platform does most of the work on its own:

- It **watches Taxscan's articles** as they're published.
- It **automatically sends the important ones** (e.g. Supreme Court & High Court rulings, ICAI, CBDT, CBIC, DGFT, IBBI, GSTAT, PMLA updates).
- It **holds back the less important ones** and **spaces sends out** so readers aren't overwhelmed.
- It **asks you to decide** on the few articles it isn't sure about (analysis/explainer pieces).

So your role is simple: **review and supervise**, and **send things manually** when you want to. This guide shows you how.

2 Logging in & your account

1. Go to `https://push.taxscan.in/admin/` in any browser.
2. Enter your **email** and **password**, then click *Log in*.
3. If this is your first time, you'll be given a temporary password and asked to set your own. Pick something strong and private.

Two kinds of account

Role	What they can do
Administrator	Everything below, <i>plus</i> add/remove team members and change roles (the Users screen).
Publisher	Use the Dashboard, Review queue, Compose/send, Campaigns and Activity. Cannot manage other users.

Tip — changing your password

You can change your password any time from the **Change password** link in the top menu.

Keep your login private

Every action you take (sending, approving, rejecting) is recorded against your name in the **Activity** log. Don't share your account.

3 The big picture: how notifications go out

Understanding this makes everything else obvious. There are **three things** the system does with every new article:

① It decides whether the article is worth a notification

Category	What happens
SEND Major authorities Supreme Court, High Court, ICAI, CBDT, CBIC, DGFT, IBBI, GSTAT, PMLA	Queued to be sent automatically.
HOLD Tribunal orders ITAT, CESTAT, NCLAT, NCLT	Not sent (these were never pushed historically). Used only as filler if there's genuinely nothing else to send.
REVIEW Everything else Analysis & explainer articles	Sent to the Review queue for <i>you</i> to approve, push, or reject.

② It paces the sending so readers aren't spammed

- **At most one push every ~45 minutes.**
- **Up to ~20 notifications per day** in total.
- **Quiet hours overnight** (roughly 11 pm–7 am) — nothing sends while people sleep.
- When several articles are waiting, it sends the **most important one first** (Supreme Court > High Court > regulatory), and the others **wait their turn** — they are never lost.

③ It lets you override, instantly, whenever you want

Automatic pacing is for the routine flow. For anything urgent, you are always in control — use **Push now** in the Review queue, or **Force send** in Compose (both explained later). These go out immediately to everyone.

Why an article didn't go out the second it was published

This is normal and intended — the system spaces sends ~45 minutes apart to protect readers. The article isn't lost; it's waiting its turn. If it's urgent, send it yourself with **Push now** / **Force send**.

4 Your daily routine — the 5-minute check

You don't need to babysit the system. A short daily glance is enough:

1. Open the Dashboard. Look at the coloured dots next to each number. ● **green** = healthy. If everything's green, you're done with this part.

2. Open the Review queue. Decide on anything waiting there — *Approve* the good analysis pieces, *Push now* the urgent ones, *Reject* the rest. Aim to leave it empty.

3. (Optional) Send anything special. Big breaking story? Use Compose to send it yourself.

That's the whole job

Glance at the Dashboard, clear the Review queue. Everything else runs automatically.

5 Reading the Dashboard

The Dashboard is your health monitor. Each metric card has a small **(i)** icon — hover it for a plain explanation. Cards that have a target show a **coloured dot**:

● **Green** — on target ● **Amber** — warning, keep an eye on it ● **Red** — below target, needs attention

What every number means & what to expect

Card	Good looks like	What it means & what to do
Active subscribers	Steadily rising	How many people are signed up to receive notifications right now — your live audience. Expect steady growth. A small sudden drop is usually just dead subscriptions being cleaned out, not lost readers.
Total sent	Always rising	The running total of notifications delivered since launch. It only ever goes up. A faster climb means more is going out and/or your audience is bigger.
Overall CTR click-through rate	● $\geq 6\%$ ● 4–6%	Of the notifications sent, the share that people actually clicked . The best gauge of how interesting and relevant your pushes are. If it's low, the headlines or the choice of what to send need work.
Opt-in rate	● $\geq 5\%$ ● 3–5%	Of the people shown the "Allow notifications?" prompt, the share who said yes . Measures how convincing the prompt is. Low usually means it appears at an awkward moment.
Delivery rate	● $\geq 95\%$ ● 90–95%	Of the notifications we tried to send, the share that actually reached the device . The rest bounce off dead/expired subscriptions. A dip usually self-heals as dead ones are removed.
Unsubscribe rate	● $< 0.5\%$	The share of subscribers who opted out or whose subscription died. Lower is better . A spike is a fatigue warning — usually too many notifications or content people didn't want.

The lower sections

- **Subscriber sources** — where your audience came from. *Recaptured* (returning visitors who already allowed notifications) is normally the biggest and should keep climbing. *Soft prompt* = fresh opt-ins from the on-site prompt.
- **30-day subscriber growth** — a simple line of new sign-ups per day. Rising or steady = healthy; flat near zero = sign-ups have stalled.
- **Opt-in funnel** — the journey from "prompt shown" → "accepted" → "subscribed". A big drop between two steps shows where readers are being lost.
- **Recent campaigns** — your latest notifications and how each performed (who it reached, clicks, CTR). Great for spotting which headlines land best.

Don't over-react to one number

Rates need a reasonable volume to be meaningful. One slow day, or a tiny send, can swing a percentage. Look at the trend over several days, not a single reading.

6 The Review queue — your main task

Open the **Review** tab daily. It lists the articles the system wasn't sure about — mostly **analysis and explainer pieces** that don't name a specific court, plus **job/recruitment posts** (see below). For each one you have three choices:

Button	What it does	Use it when...
Approve	Adds it to the normal send queue. It goes out on the next available slot, alongside everything else.	It's a good piece worth sending, but not urgent.
Push now	Sends it immediately to everyone , skipping the wait.	It's important and you want it out right away.
Reject	Drops it — it will never be sent.	It's not worth a notification.

Goal: leave the queue empty

Clearing it daily keeps the channel fresh and means no good article sits unsent. There's no penalty for rejecting — when in doubt, reject; you're curating quality.

Good to know

Tribunal orders (ITAT, CESTAT, NCLAT, NCLT) do **not** appear here — they're held automatically and need no action from you.

Job postings

Job and recruitment posts from the **Job Scan** section (CA / CMA / MBA vacancies, hiring, walk-ins, internships) now land in this Review tab automatically. **The system never sends a job post on its own — it's always your call.** Use the same three buttons: *Push now* sends it to everyone straight away, *Approve* schedules it for the next slot, *Reject* drops it.

Handling job posts

- Job posts go to **all subscribers** (there's no separate "jobs" audience), so push only the ones genuinely worth interrupting everyone for.
- Even a post that names a body (e.g. "ICAI Recruitment") comes here rather than going out automatically — that's intentional.
- **First time only:** you'll see a one-time batch of current vacancies, because the feed was just switched on. Work through them once (reject the ones not worth a push); after that only new postings appear.

Repeat headlines

If the same headline appears again **within 72 hours**, the second copy is held here instead of being pushed. When the desk publishes the same story twice as two separate posts, the system can't tell it from genuinely new news — so rather than send your readers the same headline twice, it asks you.

What to do with one

- **Genuinely the same story again?** *Reject* it — subscribers already got it.
- **A headline that legitimately recurs** — the weekly Round-Up, a digest — is just *Approve*, one click. Nothing is ever discarded automatically.
- Worth the moment it takes: getting the same headline twice is the most common reason people unsubscribe.

7 Sending a notification yourself (Compose)

Use the **Compose** screen to write and send a notification by hand — for a big story, an announcement, or anything you want out on your own terms.

The fields

Field	What to put
Title	The headline the reader sees. Keep it short and punchy — this is what makes them click.
Body	One line of detail under the title.
Click URL	The page that opens when they tap the notification — usually a taxscan.in article, but you can also link a course (academy.taxscan.in) or a product (shop.taxscan.in).
Icon URL (optional)	Leave blank to use the Taxscan logo.
Target	All subscribers , or By topic (GST, Income Tax, Customs, Corporate) to reach only people interested in that area.

Choosing who receives it — check this every time

- Compose opens on **All subscribers**. Leave it there for anything meant for the whole base (results, circulars, major rulings) — that's the normal choice.
- Switch to **By topic** only when you *deliberately* want a slice. Most subscribers picked specific topics, so a single-topic push reaches only that slice (plus the few "All news" readers) — it can land on a **very small number** of people.
- If the reach shown after sending looks far below your usual full base (~all subscribers), the target was probably narrowed by accident — just re-send with **All subscribers** selected. A duplicate only means a few people get it twice, which is harmless.
- To preview without sending to anyone, use **Test on this device** (below) — it goes only to your own browser, never to subscribers.

The two important switches

Switch	What it does & when to use it
Breaking	Lets the notification go out even during quiet hours (overnight). Use only for genuinely urgent news. It does <i>not</i> change who receives it.
Force send (full reach)	Reaches every eligible subscriber right now , ignoring the normal spacing and frequency limits. Use for a major story you want everyone to get immediately. (Without this, a hand-send still respects the usual pacing.)

Scheduling

- **Send now** or **Schedule for later** (pick a date & time). Note: Force send is immediate only.

Test on this device (preview before you send)

Below the compose box is a **Test on this device** panel — the easy way to see exactly how a notification will look before sending it to anyone.

- Click **Enable test notifications on this device** once, and accept the browser's permission prompt.

- Then click **Send test to my device** — the notification you composed appears on *your* browser only. **No one else receives it.**
- If notifications are blocked, allow them for this site in your browser settings and reload the page.

Always preview the important ones

Send a test to your own device first for anything that matters — a course or product link, a webinar, or any major push — so you can check the title, wording and that the link opens the right page.

Before you hit send

Double-check the **Title** (no typos — readers see it instantly) and the **Click URL** (open it once to be sure it's the right page). A notification can't be "unsent".

Good push, in 3 rules

Short, specific title · link straight to the article · send only what's genuinely worth interrupting someone for.

8 The other screens

Screen	What it's for
Queue	Articles already approved to send, waiting their automatic slot — listed in the order they'll go out. Click Push now on any one to send it immediately instead of waiting its turn.
Campaigns	The full history of notifications sent, each with its reach, clicks, click-rate and how much the article was actually read . Sort by any column (e.g. Pushed to see the most recent sends). The Source column shows Manual (someone sent it by hand — Compose or <i>Push now</i> — with their name) or Automatic (the system sent it on the normal schedule), so a hand-pushed item is never mistaken for an automatic one. See the column guide below.
Reports	The coverage report (what we published, by category and by court), the Reads report (what readers actually opened) and the Surfaces report (what Google Discover picked up). See below.
Activity	A record of who did what and when (logins, sends, approvals, user changes). Useful for accountability and tracing "who sent that?".
Users (Admins only)	Add teammates by email invite or with a temporary password, change their role, or deactivate them. New members get a link to set their own password.
Change password	Update your own password at any time.

Adding a teammate (Admins)

On the **Users** screen choose *Invite user* (emails them a sign-up link) or *Create user* (gives you a temporary password to share). Pick **Publisher** for editors, **Administrator** only for people who also need to manage the team.

The Campaigns screen — the columns explained

Every column is sortable — click the heading. The four that tell you how a notification performed:

Column	What it counts
Sent	How many devices we delivered the notification to.
Clicked / CTR	How many people <i>tapped the notification</i> , and that as a share of Sent. This measures the notification — was the headline worth interrupting someone for?
Reads	How many times the article page was opened, by <i>everyone</i> — Google search, the homepage, social, newsletters, and push. This measures the article, not the notification. Counts come from Google Analytics.
via Push	The slice of Reads that arrived through a push notification. Compare it against Reads to see how much of an article's traffic we generated.

Reads vs Clicked — they measure different things

Clicked is about the notification; **Reads** is about the article. Reads is normally far larger than Clicked, because most readers arrive from Google rather than from us — that is expected and not a problem. A story with high Reads but low Clicked usually means the article is strong but the push headline undersold it.

A dash (—) is not a zero

— means *no read data for this link*, not "nobody read it". You'll see it on a course or product link (academy / shop pages aren't tracked here) and on articles older than read tracking. Numbers for the last day or two also keep rising for about 48 hours as Google Analytics settles.

The Reports screen — coverage at a glance

The **Reports** tab shows two heat-maps for a period — **courts/benches × dates** (ordered Supreme Court first, then the priority High Courts, then the rest) and **categories × dates** — with the count in each cell shaded red (none) to green (plenty), plus a summary line (total articles, change vs the previous period, and where coverage was missing). It answers "*did we publish enough, across enough of the map?*"

Five tabs across the top:

Tab	What it shows
Weekly	The last 7 complete days.
Monthly	The previous calendar month.
Custom	Any window you choose — pick a <i>From</i> and <i>To</i> date and click Apply . Up to 30 days , both ends included; "To" can be today if you want a so-far-this-week look.
Reads	What readers actually opened, rather than what we published. See below.
Surfaces	What Google Discover and Google News picked up — the stories Google put in front of readers who weren't looking for us. See below.

- **Share it** — *Download image* or *Copy image* turns the whole report into a picture you can drop straight into a WhatsApp group.
- **It also arrives by email automatically** — every **Monday morning** (the past week) and on the **1st of each month** (the previous month). The email now ends with a "**How it was read**" section: reads by category and the ten most-read articles of the period.
- **Email me a test** sends the current report to your own inbox so you can see how it looks. Test emails always send the standing Weekly/Monthly report, so the button is off on the Custom, Reads and Surfaces tabs.

Who gets the report email (Admins)

Everyone with a login receives it. To include other internal colleagues who don't need a login, add their email under **Email recipients** on the Reports screen — they'll get the weekly & monthly email only, nothing else. This report is **internal**: it never goes to your notification subscribers.

The Reads tab — what readers actually opened

The coverage report answers "*what did we publish?*" The **Reads** tab answers the other half: "*what did people actually read?*" It uses the same two maps — **courts/benches** and **categories** — but instead of dates, the columns are trailing windows: **1 week · 1 month · 3 months · 6 months · 12 months**. Each cell is the number of article pageviews for that row in that window, shaded by how big a share of the period it was; hover a cell for the exact figure and how many articles it covers.

Reading it together with the coverage report is where the value is:

- **Heavy coverage, light reading** — a bench or category you publish a lot of but that barely registers in Reads. Worth asking whether it deserves the desk time.
- **Light coverage, heavy reading** — the opposite, and the more interesting signal: readers want more of it than you're producing.

- The tiles along the top give total article reads per window, so you can see the direction of travel at a glance.

Where the numbers come from

Reads are Google Analytics pageviews for taxscan.in article pages — **all** traffic, not just push. The report is rebuilt once a day, so the "data as of" line under the title may be a few hours old; that's normal and nothing needs doing.

The last two days always look low

Google Analytics keeps settling recent figures for about 48 hours, so the newest days undercount slightly and then catch up. Judge trends on the longer windows, not on yesterday.

The Surfaces tab — what Google Discover picked up

Google Discover is the feed of articles Google shows people on their phone — in the Google app and on the Chrome home screen — *without them searching for anything*. Google decides what to put there. When it picks one of our stories, a lot of readers arrive at once.

This matters more than it sounds: **Discover brings taxscan.in several times more readers than Google search does**. It simply never showed up in any report before, because Google's own analytics can't tell Discover apart from ordinary search. The Surfaces tab is the only place you can see it.

It uses the same two maps as the other tabs — **courts/benches** and **categories** — plus a list of the **most picked-up articles**. **Google News** is shown alongside, but it is tiny by comparison; if it's a small share it stays folded away behind *Show breakdown* so it doesn't crowd out the part that matters.

The columns are calendar months, oldest on the left, so you can see straight away whether a subject is rising or falling. Each row also carries a small **trend line** before the numbers — the shape of that row across every month on screen — and a **Δ** column at the far right comparing the last complete month with the one before it. The bottom row of each grid totals the whole surface.

To look at one particular stretch of time instead, use **Pick a date range** above the report: it swaps the months for the period you choose next to the *same length of time immediately before it* — so you can ask "how did the Budget fortnight compare with the fortnight before". Ranges run up to 400 days, and **Back to months** returns to the month-by-month view.

These are clicks, not Reads — don't mix them up

A number here means someone **tapped our headline in the Discover feed**. The **Reads** tab counts page views from every source. They measure different things and are collected differently, so never add them together or compare one against the other. Hover any cell for the extra detail (how often the headline was shown, and the click rate).

The newest 2–3 days are missing, not low

Google reports this data a few days late, and for the very newest days it sends *nothing at all* — not smaller numbers, none. A quiet-looking recent stretch is the delay, not a drop. That's why every column looks backwards over a period instead of showing day by day.

The last column is always short

The current month is marked with a dot and shaded with fine stripes, because Google has not finished reporting it. Every headline figure and every **Δ** therefore compares the **last complete month**, never the part-month — otherwise the report would show a collapse on the 1st of every month. History currently starts in **January 2026**; if a line appears saying history starts later, earlier months were simply never collected and their absence says nothing about what Google picked up then.

Why some changes show a count instead of a percentage

A row that went from 9 clicks to 2,800 is a true "+31,000%", which tells you nothing and swamps every real figure next to it. Where the earlier month is too small for a percentage to mean anything, the Δ shows the movement itself — **▲ 2.8k** — instead. Hover it for the exact before-and-after.

What to look for

Discover concentrates hard: a handful of stories usually account for most of the clicks. So the useful question isn't "how do we lift the average" — it's **what did those few have in common**. Look at the most-picked-up list and check whether that kind of story is something you're publishing often, or by accident.

9 Quick reference card

Daily checklist

- ☐ Dashboard — all dots green? (note anything amber/red)
- ☐ Review queue — cleared? (Approve / Push now / Reject)
- ☐ Anything urgent to send by hand?

Target numbers at a glance

Metric	● Good	● Warning	● Below target
Click-through rate (CTR)	≥ 6%	4–6%	< 4%
Opt-in rate	≥ 5%	3–5%	< 3%
Delivery rate	≥ 95%	90–95%	< 90%
Unsubscribe rate	< 0.5%	0.5–1%	> 1%

If you see this → do this

Situation	What to do
An article didn't send the instant it published	Normal — it's waiting its ~45-min turn. If urgent, Push now / Force send .
CTR is drifting down	Be choosier about what's sent; sharpen headlines. Don't send more — send better.
Unsubscribe rate ticking up	A fatigue signal — ease off volume; make sure only genuinely useful items go out.
Delivery rate dipped	Usually dead subscriptions; it self-heals. Only worry if it stays low for days.
A big story breaks	Compose → write it → Force send (add Breaking if overnight).
Review queue is piling up	Clear it: Approve the good pieces, Reject the rest.
The same headline is sitting in Review	It was published twice within 72h. Reject a true duplicate; Approve a legitimate repeat (weekly Round-Up).
A campaign shows — under Reads	No read data for that link (course/product page, or an older article) — not a zero. Nothing to fix.
Discover clicks falling month after month	Check it on complete months (the striped last column is unfinished). If the <i>All Google Discover clicks</i> total is falling too, it's real — raise it with whoever owns the site's search performance and bring the Surfaces tab as evidence.
High Reads but low Clicked	The article works; the push headline undersold it. Sharpen the wording next time.

10 Frequently asked questions

Why didn't my article go out immediately?

To protect readers, the system sends about one notification every 45 minutes. Your article is in the queue and will go out on its turn — it isn't lost. For urgent items, use **Push now** or **Force send**.

Why aren't ITAT / CESTAT / NCLAT / NCLT stories being sent?

By long-standing editorial policy these tribunal orders aren't pushed — they'd dilute the channel. The system applies this automatically. If you ever want to send one, you can do it by hand in Compose.

How do I send something to everyone, right now?

In **Compose**, make sure **Target** is set to **All subscribers** (the default), write it and tick **Force send**. Or, for an item already in the Review queue, click **Push now**.

A reader says they didn't receive a notification.

Common reasons: they never subscribed (or unsubscribed), their browser blocked notifications at the OS level, or their subscription expired (counted in the delivery rate). This is normal — no system reaches 100%.

Can I undo a send?

No — once a notification is out, it's out. Always double-check the title and link before sending. This is also why a **deleted article** is a problem: the link inside a delivered notification can't be changed, so if the page is removed from the site afterwards, anyone tapping it lands on a "page not found". If a duplicate post has to be removed, ask the web team to **redirect** it to the copy you're keeping rather than deleting it outright.

What's the difference between Clicked and Reads on the Campaigns screen?

Clicked counts people who tapped *our notification*. **Reads** counts everyone who opened *the article*, however they got there — mostly Google. Reads being much bigger than Clicked is normal. **via Push** is the overlap: the readers we brought in.

Why does a campaign show "—" in the Reads column?

It means there's no read data for that link, *not* that nobody read it. Course and product links (academy / shop) aren't tracked, and articles from before read tracking started have no figures. Very recent sends also fill in over about 48 hours.

What is Google Discover, and why is it on the Reports screen?

Discover is the feed of articles Google shows people on their phones without them searching — in the Google app and on the Chrome home screen. Google chooses what appears there. It sends taxscan.in several times more readers than Google search does, and the **Surfaces** tab is the only place it's visible, because ordinary analytics can't separate Discover traffic from normal search traffic.

Discover clicks and Reads don't match. Which is right?

Both — they count different things. A Discover **click** is someone tapping our headline in the Google feed. A **Read** is a page view of the article from any source at all. They also arrive on different schedules: Reads appear within hours, Discover figures are 2–3 days behind. Use each for its own question and don't add them together.

The Surfaces tab shows Discover falling month after month. Is the report broken?

Almost certainly not. Two things to rule out first: the **last column is always short** (it is marked with a dot and fine stripes, because Google hasn't finished reporting it), and a single quiet month can be noise. Judge it on **complete** months and on the small **trend line** at the start of each row.

If the fall shows up across most rows and in the *All Google Discover clicks* total at the foot of the grid, it is real. Discover is Google's choice, not ours — it can drop for reasons that have nothing to do with the quality of the desk's work. That makes

it a conversation to have with whoever looks after the site's search performance, and the Surfaces tab is the evidence to bring: it shows when the change started and which categories and benches it hit hardest.

Can I look at one particular stretch of time?

Yes — **Pick a date range** above the report. Choose any *From* and *To* (up to 400 days) and the grid switches to that single period alongside the *same length of time immediately before it*, with the change between them. That is the one to use for "how did the Budget fortnight compare with the fortnight before". **Back to months** returns to the month-by-month view.

The same headline turned up in my Review queue — why?

Because that headline was already published within the last 72 hours. The desk sometimes posts the same story twice, and rather than push it to your readers a second time the system asks you first. Reject it if it's a genuine duplicate; approve it if it's something that legitimately recurs, like the weekly Round-Up.

An article vanished from the queue without being sent.

The system checks each article is still live on the site immediately before pushing it. If it has been deleted in the meantime, it's archived instead of sent — that's deliberate, so a notification never points at a missing page. Nothing is wrong and no action is needed.

What's the difference between "Breaking" and "Force send"?

Breaking only lets a send happen during overnight quiet hours. **Force send** makes it reach everyone immediately, ignoring the normal spacing/frequency limits. They're independent — you can use either or both.

How often should I check in?

Once a day is plenty: a quick Dashboard glance and clearing the Review queue. The system handles the rest.

Need help?

For anything this guide doesn't cover — a number that looks wrong, a screen behaving oddly, or a question about strategy — reach out to your administrator.